

Prioritized Experience Replay

2022020425 최선묵

Dept. of Mathematics
Korea University

Experience Replay

Experience Replay (or Replay Buffer) is used to address the following two issues:

- **strongly correlated updates** that break the i.i.d. assumption of stochastic gradient algorithms.
- **the rapid forgetting of possibly rare experiences** that would be useful later on.

Using a replay memory leads to design choices at two levels:

- **which transitions to store**; and
- **which transitions to replay** (and how to do so).

Original experience replay stores **all transitions**, and they are replayed by **sampling uniformly**.

Prioritized Experience Replay addresses **the latter** by prioritizing which transitions are replayed. The priority is measured by the magnitude of their **temporal-difference (TD) error**.

Prioritized Experience Replay (PER)

- **Prioritized Experience Replay** is a way to sample the transition i with probability $P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$, measured by their **TD error**, so as to replay **important transitions** more frequently.
 - Before achieving an optimal action-value function $Q^*(s, a)$, TD-error of a transition will indicate **how ‘surprising’ or unexpected the transition is**: specifically, how far the value is from its next-step bootstrap estimate.
 - For **TD-error** $\delta_i = r_{t+1} + \gamma \max_a \hat{Q}(s_{t+1}, a; \hat{\theta}) - Q(s_t, a_t; \theta)$, the priority p_i is measured as
 - * $p_i = |\delta_i| + \epsilon$ in **proportional prioritization** (ϵ -soft: non-zero probability) ;
 - * $p_i = \frac{1}{rank(i)}$ in **rank-based prioritization**, where
 $rank(i)$ is the rank of transition i when the replay memory is sorted according to $|\delta_i|$.
It seems to be more robust as it is insensitive to outliers.
Its heavy-tail property seems to guarantee that samples will be diverse.
- In practice, however, both variants perform similarly.
- The exponent α determines how much prioritization is used ($\alpha = 0$: uniform case).

Annealing the Bias

- The prioritization **introduces bias** due to the change of sampling distribution.
- It would change the solution that the estimates will converge to.

We can correct this bias by using **importance-sampling (IS) weights** $w_i = \left(\frac{1}{N} \cdot \frac{1}{P(i)} \right)^\beta$ that fully compensates for the non-uniform probabilities $P(i)$ if $\beta = 1$.

- **Q-learning update uses $w_i \delta_i$** instead of δ_i (*weighted* IS, not ordinary IS).
For stability reasons, the weights are normalized by $1 / \max_i w_i$ (only scale downwards).
- To ensure the unbiasedness of updates at the end of training,
we define a **schedule on β that reaches 1 only at the end of learning** (e.g. linear scheduler).

Importance sampling

If samples $\{x_n\}$ are drawn from $p_X(\cdot)$, then the true mean is estimated by $\mathbb{E}_p[f(X)] \approx \sum_n f(x_n)$.

Instead of sampling from $p_X(\cdot)$, if they are sampled from another distribution $q_X(\cdot)$, we can estimate $\mathbb{E}_p[f(X)]$ by

$$\mathbb{E}_p[f(X)] = \int_X f(x) p_X(x) dx = \int_X f(x) \frac{p_X(x)}{q_X(x)} q_X(x) dx \approx \sum_n \frac{p_X(x_n)}{q_X(x_n)} f(x_n) = \sum_n w_n f(x_n) \quad \text{where } w_n = \frac{p_X(x_n)}{q_X(x_n)}.$$

Double DQN with proportional prioritization

Hyperparameters: minibatch B , replay period K and size N , exponents α and β

Initialize replay buffer \mathcal{R} to capacity N , $\Delta = 0$, $p_1 = 1$

Observe state s_1

for $t = 1, T$ **do**

 Select action $a_t = \arg \max_a Q(s_t, a; \theta)$

ϵ -greedy CNN θ

 Execute a_t in emulator and observe reward r_{t+1} and state s_{t+1}

 Store transition $(s_t, a_t, r_{t+1}, s_{t+1})$ in \mathcal{R} with maximal priority $p_t = \max_{i < t} p_i$

if $t \equiv 0 \bmod K$ **then**

for $i = 1, B$ **do**

 Sample transition $i \sim P(i) = p_i^\alpha / \sum_k p_k^\alpha$

prioritized replay + stochastic sampling

 Compute $w_i = (N \cdot P(i))^{-\beta} / \max_k w_k$

importance sampling

 Compute TD error $\delta_i = r_{i+1} + \gamma \hat{Q}(s_{i+1}, \arg \max_a Q(s_{i+1}, a; \theta); \hat{\theta}) - Q(s_i, a_i; \theta)$

CNN $\hat{\theta}$

 Update transition priority $p_i \leftarrow |\delta_i|$

 Accumulate weight-change $\Delta \leftarrow \Delta + w_i \delta_i \nabla_\theta Q(s_i, a_i; \theta)$

Double DQN

end

 Update behavior network weights $\theta \leftarrow \theta + \eta \Delta$ and reset $\Delta = 0$

 Every C steps, update target network weights $\hat{\theta} \leftarrow \theta$

end

end

Extensions

Prioritized Supervised Learning

- Analogous to PER, we can sample each data using a priority based on its last-seen error.
- This can help focus the learning on those samples that can still be learned from.

Off-policy Replay

- Apply the replay probability and the IS-correction to importance sampling or rejection sampling.

Feedback for Exploration

- M_i : the total number that the i th transition will end up being replayed.
 - * This varies widely, and this gives a rough indication of how useful it was to the agent.
- We can monitor usefulness of the experience via M_i and update the distribution toward generating more useful experience (feedback to exploration strategy).


Prioritized Memories

- Priority can help determine which transitions to store and when to erase them.

How to implement Prioritized Experience Replay?

Proportional Prioritization

- An efficient implementation can be made using a 'sum-tree' data structure.
- The value written on a leaf node is the frequency that each leaf node is randomly sampled at.

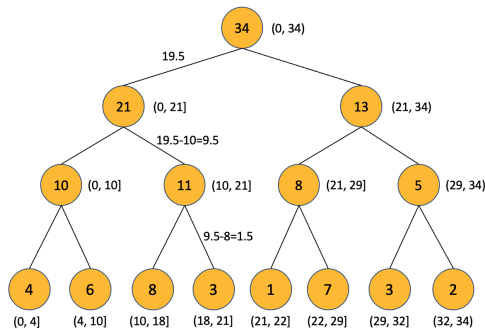


A **binary tree** is a tree data structure where each parent node has a maximum of 2 child nodes.

A **sum tree** is a binary tree where every node is the sum of its children.

PER can be implemented as a sum tree with the priorities in the leaf nodes.

How to sample a leaf node in Sum Tree



V = the value of the root node
Sample a number x from $Unif(0, V)$.

repeat until we reach a leaf node.

v_1, v_2 = the value of left and right child nodes, resp.

if $x < v_1$

choose the left child node.

else choose the right child node

$x = x - v_1$


end

return the chosen node

- We split the interval $(0, V)$ into the sub-intervals of the length of the values in the leaf nodes.
Sampling a leaf node = finding a sub-interval that contains the sampled number x from $Unif(0, V)$.
- To sample a minibatch of size B , the range $(0, V)$ is divided equally into B ranges.
A value is uniformly sampled from each range to sample a minibatch.

How to Update the Value of a leaf node in Sum Tree

- Using a sum tree structure, the value of a leaf node can be updated by the following algorithm.
- Notice that the algorithm has a complexity of $O(\log_2 N)$, instead of $O(N)$.



s, v : a leaf node and its value.

Δ : the change of value of a leaf node s .

$v = v + \Delta$

repeat until we reach the root node.

if there is a parent node s' of s ,

$v' = v' + \Delta$ (v' : the value of s')

$s = s'$

end