

Proximal Policy Optimization

Sunmook Choi

Dept. of Mathematics
Korea University

Proximal Policy Optimization

TRPO Goal: $\max_{\theta} \hat{\mathbb{E}}_t \left[\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} \hat{A}_t \right]$ subject to $\hat{\mathbb{E}}_t \left[\text{KL}(\pi_{\theta_{\text{old}}}(\cdot|s_t) \parallel \pi_{\theta}(\cdot|s_t)) \right] \leq \delta$

However, TRPO is **too complicated to implement** and **too heavy to compute**:

- * TRPO algorithm contains 2nd-order optimization (natural policy gradient).

The motivation of PPO is the same as that of TRPO.

- How can we update the policy **as big as possible**?
- But we want to update **not too much** so that we can prevent an accidental disaster.
- At the same time, we want to make it easy to implement and to have less computation.

Clipped Surrogate Objective

$$L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)} \hat{A}_t, \text{clip} \left(\frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}, 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right) \right]$$

- \hat{A}_t : an advantage-function estimator

Notice that, for $r_t(\theta) = \frac{\pi_{\theta}(a_t|s_t)}{\pi_{\theta_{\text{old}}}(a_t|s_t)}$, the update equation is equal to the following:

- * For $\hat{A}_t > 0$, $L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, (1 + \epsilon) \hat{A}_t \right) \right]$: *Encouraging* the chosen action a_t .
- * For $\hat{A}_t < 0$, $L^{CLIP}(\theta) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta) \hat{A}_t, (1 - \epsilon) \hat{A}_t \right) \right]$: *Discouraging* the chosen action a_t .
- * The clipping value $\epsilon > 0$ is a hyperparameter.

Adaptive KL Penalty Coefficient

The authors provide an alternative to the clipped surrogate objective, or in addition to it. The approach uses a **penalty on KL divergence** and **adapts the penalty coefficient**.

- Using several epochs of minibatch SGD, optimize the KL-penalized objective

$$L^{KL PEN}(\theta) = \hat{\mathbb{E}}_t \left[r_t(\theta) \hat{A}_t - \beta \text{KL}(\pi_{\theta_{\text{old}}}(\cdot|s_t) \parallel \pi_{\theta}(\cdot|s_t)) \right]$$

- Compute $d = \hat{\mathbb{E}}_t \left[\text{KL}(\pi_{\theta_{\text{old}}}(\cdot|s_t) \parallel \pi_{\theta}(\cdot|s_t)) \right]$
 - * If $d < d_{\text{targ}}/1.5$, then $\beta \leftarrow \beta/2$. (i.e., **encouraging** policy update)
 - * If $d > d_{\text{targ}} \times 1.5$, then $\beta \leftarrow \beta \times 2$. (i.e., **discouraging** policy update)

The hyperparameter d_{targ} denotes the target of the amount of changes at each policy update.

Advantage Function Estimator

- One style of estimating the advantage function is to run the policy for fixed T timesteps:

$$\hat{A}_t = -V(s_t) + r_t + \gamma r_{t+1} + \cdots + \gamma^{T-t+1} r_{T-1} + \gamma^{T-t} V(s_T)$$

which is used in the A3C paper.

- Another style is to use a truncated version of **generalized advantage estimation (GAE)**, a generalization of the choice above, which reduces to the one above when $\lambda = 1$:

$$\hat{A}_t = \delta_t + (\gamma\lambda)\delta_{t+1} + \cdots + (\gamma\lambda)^{T-t+1}\delta_{T-1},$$

$$\text{where } \delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$$

Recall: TD(λ) (오승상 교수님 강화학습 강의자료 p.62)

- n -step return
 - * $G_t^{(1)} = R_{t+1} + \gamma V(S_{t+1})$
 - * $G_t^{(n)} = R_{t+1} + \gamma R_{t+2} + \cdots + \gamma^{n-1} R_{t+n} + \gamma^n V(S_{t+n})$
- $G_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} G_t^{(n)}$ for $\lambda \in [0, 1]$.
 - * the exponentially-weighted average of n -step returns $G_t^{(n)}$
 - * TD(λ) is a method of updating value function as follows:

$$V(S_t) \leftarrow V(S_t) + \alpha[G_t^\lambda - V(S_t)]$$

- * If $\lambda = 0$, then it is equal to the original TD update equation.

Generalized Advantage Estimation

With the definition $\delta_t = r_t + \gamma V(s_{t+1}) - V(s_t)$,

$$* \hat{A}_t^{(1)} = \delta_t = -V(s_t) + r_t + \gamma V(s_{t+1})$$

$$* \hat{A}_t^{(k)} = \sum_{l=0}^{k-1} \gamma^l \delta_{t+l} = -V(s_t) + r_t + \gamma r_{t+1} + \dots + \gamma^{k-1} r_{t+k-1} + \gamma^k V(s_{t+k})$$

$$* \hat{A}_t^{(\infty)} = \sum_{l=0}^{\infty} \gamma^l \delta_{t+l} = -V(s_t) + \sum_{l=0}^{\infty} \gamma^l r_{t+l} \quad \longleftarrow \quad \text{an advantage-function estimator.}$$

$\hat{A}_t^{\text{GAE}(\gamma, \lambda)}$: the generalized advantage estimator

* the exponentially-weighted average of these k -step estimators

$$* \hat{A}_t^{\text{GAE}(\gamma, \lambda)} = (1 - \lambda) \left(\hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots \right) = \dots = \sum_{l=0}^{\infty} (\gamma \lambda)^l \delta_{t+l}$$

Practical Implementation

Final objective: $L_t^{CLIP+VF+S}(\theta) = \hat{\mathbb{E}}_t \left[L_t^{CLIP}(\theta) - c_1 L_t^{VF}(\theta) + c_2 S[\pi_\theta](s_t) \right]$

- $L_t^{VF}(\theta) = (V_\theta(s_t) - V_t^{targ})^2$: to learn the **state-value function**
 - * $V_t^{targ} = r_{t+1} + \gamma V_\theta(s_{t+1})$: TD-target
 - * $V_t^{targ} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{T-1} r_T = G_t$: MC-target
- $S[\pi_\theta](s_t) = - \sum_{a \in \mathcal{A}} \pi_\theta(a|s_t) \log \pi_\theta(a|s_t)$: **entropy bonus** to ensure exploration